

CSE3211: Operating System Assignment 0

Al Md. Aladin

student ID: 2015-516-820 & roll number: FH-59

Sharear Bin Amin

student ID: 2015-016-816 & roll number: FH-55

August 13, 2018

1 Questions & Answers

Q1. What is the vm system called that is configured for assignment 0?

Answer: `dumbvm` from `kern/arch/mips/conf/conf.arch`

Q2. Which register number is used for the stack pointer (sp) in OS/161?

Answer: `#define sp $29 /* stack pointer */` from `kern/arch/mips/include/kern/regdefs.h`

Q3. What bus/busses does OS/161 support?

Answer: `LAMEbus` from `kern/arch/sys161/include/bus.h`

Q4. What is the difference between `splhigh` and `spl0`?

Answer: `splhigh()` sets **IPL** to the highest value, disabling all interrupts. `spl0()` sets **IPL** to 0, enabling all interrupts. from `os161-ASST0/kern/include/spl.h`

Q5. Why do we use typedefs like `u_int32_t` instead of simply saying "int"?

Answer: To make sure that we really get a 32-bit unsigned integer (unsigned int depends on the platform) from `kern/arch/mips/include/types.h`

Q6. What must be the first thing in the process control block?

Answer: `pcb_switchstack` must be the first thing in the process control block.

Q7. What does `splx` return?

Answer: The old interrupt state (an integer) from `os161-ASST0/kern/include/spl.h`

Q8. What is the highest interrupt level?

Answer: `#define IPL_HIGH 1` from `os161-ASST0/kern/include/spl.h`

Q9. What function is called when user-level code generates a fatal fault?

Answer: `void kill_curthread(vaddr_t epc, unsigned code, vaddr_t vaddr)` from `os161-ASST0/kern/arch/mips/mips/trap.c`

Q10. How frequently are hardclock interrupts generated?

Answer: 100 times a second `#define HZ 100` from `kern/include/clock.h`

Q11. What functions comprise the standard interface to a VFS device?

Answer: `vfs_setcurdir`, `vfs_clearcurdir`, `vfs_getcurdir`, `vfs_sync`, `vfs_getroot`, `vfs_getdevname`, `vfs_lookup`, `vfs_lookupparent`, `vfs_open`, `vfs_close`, `vfs_readlink`, `vfs_symlink`, `vfs_mkdir`, `vfs_link`, `vfs_remove`, `vfs_rmdir`, `vfs_rename`, `vfs_chdir`, `vfs_getcwd`, `vfs_bootstrap`, `vfs_initbootfs`, `vfs_setbootfs`, `vfs_clearbootfs`, `vfs_adddev`, `vfs_addfs`, `vfs_mount`, `vfs_unmount`, and `vfs_unmountall` from `os161-ASST0/kern/include/vfs.h`

Q12. How many characters are allowed in a volume name?

Answer: `#define SFS_VOLNAME_SIZE 32` /* max length of volume name */ from `kern/include/kern/sfs.h`

Q13. How many direct blocks does an SFS file have?

Answer: `#define SFS_NDIRECT 15` /* # of direct blocks in inode */ from `kern/include/kern/sfs.h`

Q14. What is the standard interface to a file system (i.e., what functions must you implement to implement a new file system)?

Answer: `fsop_sync`, `fsop_getvolname`, `fsop_getroot`, `fsop_unmount` from `kern/include/fs.h`

Q15. What function puts a thread to sleep?

Answer: `Void wchan_sleep(struct wchan *wc, struct spinlock *lk)` from `kern/thread/thread.c`

Q16. How large are OS/161 pids?

Answer: `typedef int32_t pid_t;` /* Process ID */ 32 bits / 4 bytes from `kern/include/kern/types.h`

Q17. What operations can you do on a vnode?

Answer: `vop_eopen, vop_reclaim, vop_read, vop_readlink, vop_getdirentry, vop_write, vop_ioctl, vop_stat, vop_gettype, vop_tryseek, vop_fsync, vop_mmap, vop_truncate, vop_namefile, vop_creat, vop_symlink, vop_mkdir, vop_link, vop_remove, vop_rmdir, vop_rename, vop_lookup, vop_lookupparent` from `kern/include/vnode.h`

Q18. What is the maximum path length in OS/161?

Answer: * Longest full path name */ `#define PATH_MAX 1024` from `kern/include/kern/limits.h`

Q19. What is the system call number for a reboot?

Answer: The system call number for a reboot is 119.

Q20. Where is `STDIN_FILENO` defined?

Answer: `#define STDIN_FILENO 0` /* Standard input */ from `kern/include/kern/unistd.h`

Q21. What does `kmain()` do?

Answer: Kernel main. (Boot up, then fork the menu thread, wait for a reboot request, and then shut down.) from `kern/main/main.c`

Q22. Is it OK to initialise the thread system before the scheduler? Why (not)?

Answer: Yes. The scheduler bootstrap just creates the run queue, and the thread bootstrap just initializes the first thread.

Q23. What is a zombie?

Answer: “Zombies are threads/processes that have exited but not been fully deleted yet.” from `kern/thread/thread.c`

Q24. How large is the initial run queue?

Answer: The initial run queue is 32.

Q25. What does a device name in OS/161 look like?

Answer: The name of a device is always `device:`, such as `lhd0:` from `kern/vfs/device.c`

Q26. What does a raw device name in OS/161 look like?

Answer: The name with raw appended, such as `lhd0raw`, from

kern/vfs/vfslist.c

Q27. What lock protects the vnode reference count?

Answer: vn_countlock protects the vnode reference count from kern/vfs/vnode.c

Q28. What device types are currently supported?

Answer: The device types currently supported are block and character devices. from kern/vfs/device.c

.....